

students_post

March 11, 2024

0.1 Logistic Regression - Gradient Descent

In this part you will build a logistic regression model using Numpy and doing gradient descent. You should complete the following cells (those with comments and no code).

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
```

```
[ ]: # read data
import pandas as pd
from sklearn.preprocessing import add_dummy_feature

data = pd.read_csv('https://raw.githubusercontent.com/thomouvic/SENG474/main/
↳data/exams_admitted.csv')
print(data.head())
```

	exam1	exam2	admitted
0	34.623660	78.024693	0
1	30.286711	43.894998	0
2	35.847409	72.902198	0
3	60.182599	86.308552	1
4	79.032736	75.344376	1

```
[ ]: # extract X and y from data
# use features 'exam1' and 'exam2' for X
# use feature 'admitted' for y
# use pandas.DataFrame.values to convert to numpy arrays
X = data[['exam1', 'exam2']].values
y = data[['admitted']].values

print(X)
```

```
[[34.62365962 78.02469282]
 [30.28671077 43.89499752]
 [35.84740877 72.90219803]
 [60.18259939 86.3085521 ]
 [79.03273605 75.34437644]
 [45.08327748 56.31637178]
 [61.10666454 96.51142588]
 [75.02474557 46.55401354]
```

[76.0987867 87.42056972]
[84.43281996 43.53339331]
[95.86155507 38.22527806]
[75.01365839 30.60326323]
[82.30705337 76.4819633]
[69.36458876 97.71869196]
[39.53833914 76.03681085]
[53.97105215 89.20735014]
[69.07014406 52.74046973]
[67.94685548 46.67857411]
[70.66150955 92.92713789]
[76.97878373 47.57596365]
[67.37202755 42.83843832]
[89.67677575 65.79936593]
[50.53478829 48.85581153]
[34.21206098 44.2095286]
[77.92409145 68.97235999]
[62.27101367 69.95445795]
[80.19018075 44.82162893]
[93.1143888 38.80067034]
[61.83020602 50.25610789]
[38.7858038 64.99568096]
[61.37928945 72.80788731]
[85.40451939 57.05198398]
[52.10797973 63.12762377]
[52.04540477 69.43286012]
[40.23689374 71.16774802]
[54.63510555 52.21388588]
[33.91550011 98.86943574]
[64.17698887 80.90806059]
[74.78925296 41.57341523]
[34.18364003 75.23772034]
[83.90239366 56.30804622]
[51.54772027 46.85629026]
[94.44336777 65.56892161]
[82.36875376 40.61825516]
[51.04775177 45.82270146]
[62.22267576 52.06099195]
[77.19303493 70.4582]
[97.77159928 86.72782233]
[62.0730638 96.76882412]
[91.5649745 88.69629255]
[79.94481794 74.16311935]
[99.27252693 60.999031]
[90.54671411 43.39060181]
[34.52451385 60.39634246]
[50.28649612 49.80453881]
[49.58667722 59.80895099]

```

[97.64563396 68.86157272]
[32.57720017 95.59854761]
[74.24869137 69.82457123]
[71.79646206 78.45356225]
[75.39561147 85.75993667]
[35.28611282 47.02051395]
[56.2538175 39.26147251]
[30.05882245 49.59297387]
[44.66826172 66.45008615]
[66.56089447 41.09209808]
[40.45755098 97.53518549]
[49.07256322 51.88321182]
[80.27957401 92.11606081]
[66.74671857 60.99139403]
[32.72283304 43.30717306]
[64.03932042 78.03168802]
[72.34649423 96.22759297]
[60.45788574 73.0949981 ]
[58.84095622 75.85844831]
[99.8278578 72.36925193]
[47.26426911 88.475865 ]
[50.4581598 75.80985953]
[60.45555629 42.50840944]
[82.22666158 42.71987854]
[88.91389642 69.8037889 ]
[94.83450672 45.6943068 ]
[67.31925747 66.58935318]
[57.23870632 59.51428198]
[80.366756 90.9601479 ]
[68.46852179 85.5943071 ]
[42.07545454 78.844786 ]
[75.47770201 90.424539 ]
[78.63542435 96.64742717]
[52.34800399 60.76950526]
[94.09433113 77.15910509]
[90.44855097 87.50879176]
[55.48216114 35.57070347]
[74.49269242 84.84513685]
[89.84580671 45.35828361]
[83.48916274 48.3802858 ]
[42.26170081 87.10385094]
[99.31500881 68.77540947]
[55.34001756 64.93193801]
[74.775893 89.5298129 ]]

```

```

[ ]: # normalize X
      # use scikit-learn's built-in function MinMaxScaler

```

```

from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
X = scaler.fit_transform(X)
print(X[0])

```

```
[0.06542784 0.69465488]
```

```

[ ]: # add a dummy feature for the intercept
# use scikit-learn's built-in function add_dummy_feature
from sklearn.preprocessing import add_dummy_feature

X = add_dummy_feature(X)
print(X[0])

```

```
[1.          0.06542784 0.69465488]
```

```

[ ]: # set m (number of training examples) and n (number of features)
# use the shape attribute of X

m=X.shape[0]
n=X.shape[1]

m,n

```

```
[ ]: (100, 3)
```

```

[ ]: # initialize theta to zeros
# the theta array should be number of features plus one

theta = np.zeros((n,1)) # this is number of features plus one cause we dont
    ↪ count dummy i think
theta

```

```

[ ]: array([[0.],
           [0.],
           [0.]])

```

```

[ ]: # define sigmoid function
def sigmoid(z):
    return 1 / (1+np.exp(-z))

# test your sigmoid function on the value 0, should return 0.5
sigmoid(0)

```

```
[ ]: 0.5
```

```
[ ]: # create a hypothesis function called h that takes in:
# theta, an instance x, and returns the hypothesis
# the hypothesis is the sigmoid of x@theta
# use the @ operator for matrix multiplication

print(theta)

def h(theta,x):
    return sigmoid(x@theta)

print(X[0])

# test your hypothesis function on the first instance of X, should return [[0.
↪5]]
print(h(theta,X[0]))

# the above function is vectorized
# for example, if instead of a single instance x, we have a matrix X of shape ↪
↪(m,n)
# then the hypothesis is a vector of shape (m,1)
# where each element is the hypothesis for the corresponding row of X
# test it on the first 5 instances of X, should return an array of 0.5's

h(theta,X[:5])
```

```
[[0.]
 [0.]
 [0.]]
[1.          0.06542784 0.69465488]
[0.5]
```

```
[ ]: array([[0.5],
           [0.5],
           [0.5],
           [0.5],
           [0.5]])
```

```
[ ]: import numpy as np
# create a function called J that takes in theta, X, y, and returns the cost
# the cost is the average of the log loss over the training examples
# the log loss for a single example is ↪
↪-y*log(h(theta,x))-(1-y)*log(1-h(theta,x))
# the cost is the average of the log loss over the training examples
# use the np.mean function to compute the average
# use the np.log function to compute the log
# use the @ operator to compute matrix multiplication
```

```

# use a vectorized implementation, do not use a for loop over the training
  ↪ examples
# use the hypothesis function h defined above

def J(theta,X,y):
    one_c = -y*np.log(h(theta,X))-(1-y)*np.log(1-h(theta,X))
    aver = np.mean(one_c)
    #print(one_c)
    #print(aver)
    return aver

#test your cost function on the initial all-zero theta, should return 0.
  ↪ 6931471805599453
J(theta,X,y)

```

```
[ ]: 0.6931471805599453
```

```

[ ]: # create a function called gradient that takes in theta, X, y, and returns the
  ↪ gradient
# the gradient is the average of the gradient over the training examples
# use the hypothesis function h defined above
# use a vectorized implementation, do not use a for loop over the training
  ↪ examples
# use the @ operator to compute matrix multiplication
# use the np.mean function to compute the average
# use the formula for the gradient given in the lecture
# the vectorized formula is  $X.T @ (h(\theta, X) - y) / m$ 

# test your gradient function on the initial theta

def gradient(theta, X, y):
    grad = X.T @ (h(theta,X)-y)/m
    # print(grad)
    return grad

gradient(theta,X,y)

```

```
[ ]: array([[ -0.1         ],
           [-0.12904484],
           [-0.12015491]])
```

```

[ ]: # create a function called 'fit' that takes in:
# X, y, alpha, num_iters, initial_theta,
# and returns: final_theta, and J_history
# inside the function:
# initialize theta to the initial_theta
# initialize J_history to an empty list
# for each iteration, using a for loop over num_iters:
# update theta by subtracting alpha times the gradient (using the gradient_
    ↪function defined above)
# compute the cost J using the cost function defined above and append it to_
    ↪J_history
# return final_theta, and J_history

def fit(X,y,alpha,num_iters,initial_theta):
    theta = initial_theta
    J_history = []
    for q in range(0, num_iters):
        theta = theta - (alpha * gradient(theta, X, y))
        #print(theta)
        j = J(theta, X, y)
        J_history.append(j)
    final_theta = theta
    #print(final_theta, J_history)
    return final_theta, J_history

#fit(X,y,0.1,10000,np.zeros((n,1)))

# create a function called predict that takes in:
# theta, and an array X_new of instances,
# and returns the predictions
# threshold the hypothesis at 0.5
# use the hypothesis function h defined above

def predict(theta, X_new):
    predictions = h(theta, X_new[0])
    print(predictions)

[ ]: # call fit() with the following arguments:
# X, y, alpha=0.1, num_iters=10000, initial_theta=np.zeros((n,1))
# store the returned values in theta, J_history
# uncomment the following line to call fit()

theta, J_history = fit(X, y, alpha=0.1, num_iters=10000, initial_theta=np.
    ↪zeros((n,1)))

```


```

# plot the cost over the iterations stored in J_history
# you should see the cost decreasing
# uncomment the following lines to plot the cost

plt.plot(J_history)
plt.xlabel('Iteration')
plt.ylabel('J')

```

```
[ ]: Text(0, 0.5, 'J')
```



```

[ ]: # plot the data points
# use a scatter plot
# use the first feature for the x-axis, the second feature for the y-axis
# use the actual labels for the color, c=y[:,0]
# uncomment the following line to plot the data points
plt.scatter(X[:,1],X[:,2],c=y[:,0])

# plot the decision boundary

```



```

# the decision boundary is the line where the hypothesis is 0.5
# the hypothesis is 0.5 when  $x @ \theta = 0$ 
# so the decision boundary is the line where  $x @ \theta = 0$ 
# this is a line in the  $x_1, x_2$  plane

# for plotting the decision boundary, we need two points
# create two  $x_1$  values, say 0 and 1 (since we scaled to  $[0,1]$ )
# then calculate the corresponding  $x_2$  values
# using the decision boundary equation
# uncomment the following lines to plot the decision boundary
two_x1 = np.array([0, 1])
two_x2 = -(theta[0] + theta[1] * two_x1) / theta[2]

# plot the decision boundary as a  $k--$  line.  $k--$  is black dashed line
plt.plot(two_x1, two_x2, "k--", linewidth=3)

```

```
[ ]: [<matplotlib.lines.Line2D at 0x2cddfdeb5c0>]
```

