

Text_Classification

March 11, 2024

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import re
import nltk
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer
from nltk.tokenize import word_tokenize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import cross_val_score

# suppress warnings
import warnings
warnings.filterwarnings('ignore')
```

```
[ ]: # read the data
df = pd.read_csv('https://raw.githubusercontent.com/nikjohn7/
↳Disaster-Tweets-Kaggle/main/data/train.csv')
df.head()
```

```
[ ]:   id keyword location                                text \
0    1     NaN      NaN  Our Deeds are the Reason of this #earthquake M...
1    4     NaN      NaN                Forest fire near La Ronge Sask. Canada
2    5     NaN      NaN  All residents asked to 'shelter in place' are ...
3    6     NaN      NaN  13,000 people receive #wildfires evacuation or...
4    7     NaN      NaN  Just got sent this photo from Ruby #Alaska as ...

      target
0          1
1          1
2          1
3          1
4          1
```

```
[ ]: # how many rows and columns are in the data set?
df.shape
```

```
[ ]: (7613, 5)
```

```
[ ]: # how many speeches are there from the United States?
      #df[df['country'] == 'USA'].shape

      #how many targets equal to one?
      df[df['target']==1].shape
```

```
[ ]: (3271, 5)
```

```
[ ]: # how many speeches are there from CANADA?
      #df[df['country'] == 'CAN'].shape

      #how many targets equal to 10?
      df[df['target'] == 0].shape
```

```
[ ]: (4342, 5)
```

```
[ ]: # show rows for the United States
      #df[df['country'] == 'USA']

      #show rows for target 1
      df[df['target'] == 1]
```

```
[ ]:
      id keyword location \
0      1      NaN      NaN
1      4      NaN      NaN
2      5      NaN      NaN
3      6      NaN      NaN
4      7      NaN      NaN
...
7608 10869      NaN      NaN
7609 10870      NaN      NaN
7610 10871      NaN      NaN
7611 10872      NaN      NaN
7612 10873      NaN      NaN
```

```

      text target
0      Our Deeds are the Reason of this #earthquake M...      1
1      Forest fire near La Ronge Sask. Canada      1
2      All residents asked to 'shelter in place' are ...      1
3      13,000 people receive #wildfires evacuation or...      1
4      Just got sent this photo from Ruby #Alaska as ...      1
...
7608 Two giant cranes holding a bridge collapse int...      1
7609 @aria_ahrury @TheTawniest The out of control w...      1
7610 M1.94 [01:04 UTC]?5km S of Volcano Hawaii. htt...      1
7611 Police investigating after an e-bike collided ...      1
```

7612 The Latest: More Homes Razed by Northern Calif... 1

[3271 rows x 5 columns]

```
[ ]: # show rows for target 0
df[df["target"] == 0 ]
```

```
[ ]:
      id keyword location \
15     23      NaN      NaN
16     24      NaN      NaN
17     25      NaN      NaN
18     26      NaN      NaN
19     28      NaN      NaN
...
7581 10833 wrecked Lincoln
7582 10834 wrecked      NaN
7584 10837      NaN      NaN
7587 10841      NaN      NaN
7593 10848      NaN      NaN
```

		text	target
15		What's up man?	0
16		I love fruits	0
17		Summer is lovely	0
18		My car is so fast	0
19		What a goooooooooaaaaaal!!!!!!	0
...	
7581	@engineshed	Great atmosphere at the British Li...	0
7582	Cramer:	Iger's 3 words that wrecked Disney's s...	0
7584	These boxes	are ready to explode! Exploding Ki...	0
7587		Sirens everywhere!	0
7593	I just heard	a really loud bang and everyone i...	0

[4342 rows x 5 columns]

```
[ ]: import nltk          #including and excluding stopwords nltk library  #we will
      ↪probably use this later
nltk.download('stopwords')

stopwords = set(nltk.corpus.stopwords.words('english'))

include_stopwords =
    ↪{'normal','routine','everyday','regular','common','typical','usual','standard','average','g
    ↪'often'}
#exclude_stopwords = {'against'}

stopwords |= include_stopwords
```

```
#stopwords -= exclude_stopwords
```

```
[nltk_data] Downloading package stopwords to  
[nltk_data] C:\Users\naumh\AppData\Roaming\nltk_data...  
[nltk_data] Package stopwords is already up-to-date!
```

```
[ ]: # build a text processing and classifier pipeline  
# to predict the country (USA or Canada) of a speech  
  
from sklearn.feature_extraction.text import TfidfVectorizer  
from sklearn import svm  
from sklearn.model_selection import train_test_split  
from sklearn.pipeline import Pipeline  
from sklearn.metrics import classification_report  
  
#df2 = df[df['country'].isin(['USA', 'CAN'])]  
df2 = df[df['target'].isin([0,1])]  
  
# Split the dataset into training and test sets  
X_train, X_test, y_train, y_test = train_test_split(df2['text'], df2['target'],  
↳test_size=0.4)  
  
# Create a pipeline that first transforms the text data into TF-IDF vectors,  
↳then applies SVM  
text_clf = Pipeline([  
    ('tfidf', TfidfVectorizer(stop_words=list(stopwords))), #sparse vectors  
    ↳what if words not in classical vocab- words that show up in classical  
    ↳collection  
    ('clf', svm.SVC()),  
]) #two weights for each word in the document. TF means term frequency, idf  
↳means inverse document frequency, rarer word, higher it is in the document  
  
#"terms" are just words. Weighting scemet to capture how important word is in  
↳respect to document. Used in search engines, text classification, other  
↳natural language processing tasks  
  
#weight words by how often appear in document  
#how rare they are across all documents in a collection IDF  
  
#TF(t,d) = Number of times term t occurs in document d / total number of terms  
↳in document d IDF is inverse fraction of documents that contain word, whic  
↳his then scaled logarithmically  
  
#IDF(t) = log((1+n)/1+df(T)) + 1 where n is #doc. df is #occur of t in  
↳collection
```

```

#df(t) is the number of documents in the collection that contain term t

#TFIDF(t,d) = TF(t,d) * IDF(t)

#check 'toy' example in PDF

#2.1 COSINE SIMILARITY (scale invariant) based on size of vocabulary. IT'S U
    ↪ VERY IMPORTANT, different than euclidean.

#we use tf idf vectorizer in notebook

# Train the classifier
text_clf.fit(X_train, y_train)

# Predict the test set results
y_pred = text_clf.predict(X_test)

# Print the classification report
print(classification_report(y_test, y_pred, target_names=['0', '1']))

```

	precision	recall	f1-score	support
0	0.78	0.87	0.82	1749
1	0.79	0.67	0.72	1297
accuracy			0.78	3046
macro avg	0.78	0.77	0.77	3046
weighted avg	0.78	0.78	0.78	3046

```

[ ]: # This script creates a new column 'sentiment' in the dataframe,
# which contains the sentiment score of the text.
# The sentiment score is a float within the range [-1.0, 1.0],
# where -1.0 denotes a very negative sentiment,
# 1.0 denotes a very positive sentiment,
# and values around 0 denote a neutral sentiment.

from textblob import TextBlob

# Define a function to apply sentiment analysis to a text
def get_sentiment(text):
    blob = TextBlob(text)
    return blob.sentiment.polarity # returns a value between -1 and 1

```

```
# Create a new column 'sentiment' in the DataFrame
#df2['sentiment'] = df2['text'].apply(get_sentiment)
df2['sentiment'] = df2['text'].apply(get_sentiment)

# Display the DataFrame
df2
```

```
[ ]:      id keyword location \
0         1      NaN      NaN
1         4      NaN      NaN
2         5      NaN      NaN
3         6      NaN      NaN
4         7      NaN      NaN
...
7608  10869      NaN      NaN
7609  10870      NaN      NaN
7610  10871      NaN      NaN
7611  10872      NaN      NaN
7612  10873      NaN      NaN

      text  target  sentiment
0  Our Deeds are the Reason of this #earthquake M...      1  0.000000
1      Forest fire near La Ronge Sask. Canada      1  0.100000
2  All residents asked to 'shelter in place' are ...      1 -0.018750
3  13,000 people receive #wildfires evacuation or...      1  0.000000
4  Just got sent this photo from Ruby #Alaska as ...      1  0.000000
...
7608  Two giant cranes holding a bridge collapse int...      1  0.000000
7609  @aria_ahrury @TheTawniest The out of control w...      1  0.150000
7610  M1.94 [01:04 UTC]?5km S of Volcano Hawaii. htt...      1  0.000000
7611  Police investigating after an e-bike collided ...      1 -0.260417
7612  The Latest: More Homes Razed by Northern Calif...      1  0.500000

[7613 rows x 6 columns]
```

```
[ ]: # find average sentiment for each country in df2
#df2.groupby('country')['sentiment'].mean()
df2.groupby('target')['sentiment'].mean()
```

```
[ ]: target
0    0.070622
1    0.018631
Name: sentiment, dtype: float64
```

```
[ ]: # find average sentiment for each speaker in df2
# order the results from most positive to most negative
```

```
#df2.groupby('speaker')['sentiment'].mean().sort_values(ascending=False).head(5)
df2.groupby('keyword')['sentiment'].mean().sort_values(ascending=False).head(5)
```

```
[ ]: keyword
hazardous      0.457891
razed          0.418946
outbreak       0.312661
mayhem         0.277262
wreckage       0.273440
Name: sentiment, dtype: float64
```

```
[ ]: #df2.groupby('year')['sentiment'].mean().sort_values(ascending=False)
df2.groupby('location')['sentiment'].mean().sort_values(ascending=False)
```

```
[ ]: location
The Waystone Inn      1.0
Florida USA           1.0
Paranaque City        1.0
West Palm Beach, Florida 1.0
Washington, DC 20009  1.0
...
Mumbai india          -1.0
Thibodaux, LA         -1.0
fujo garbage heaven   -1.0
sri lanka              -1.0
Milton Keynes, England -1.0
Name: sentiment, Length: 3341, dtype: float64
```